

RTML-Final-2021

May 7, 2021

1 RTML Final 2021

In this exam, we'll have some practical exercises using RNNs and some short answer questions regarding the Transformer/attention and reinforcement learning.

Consider the AGNews text classification dataset:

```
[60]: from torchtext.datasets import AG_NEWS
      from torchtext.data.utils import get_tokenizer
      from collections import Counter
      from torchtext.vocab import Vocab
      import pandas as pd
      import string
      import re

      #train_iter = pd.read_csv("./data/.data/train.csv")
      #train_iter['summary'] = train_iter['Title'] + ' ' + train_iter['Description']
      train_iter = AG_NEWS(root='.data',split='train')
      tokenizer = get_tokenizer('basic_english')
      counter = Counter()

      def clean(line):
          line = line.replace('\\', ' ')
          line = re.sub('(\s+)(a|an|and|the)(\s+)', '\1\3', line)
          line = re.sub('[%s]' % re.escape(string.punctuation), '', line)

          return line

      labels = {}
      for (label, line) in train_iter:
          if label in labels:
              labels[label] += 1
          else:
              labels[label] = 1
          counter.update(tokenizer(clean(line)))

      vocab = Vocab(counter, min_freq=0, max_size=1000)
```

```

print('Label frequencies:', labels)
print('A few token frequencies:', vocab.freqs.most_common(5))
print('Label meanings: 1: World news, 2: Sports news, 3: Business news, 4: Sci/
↳Tech news')

```

Label frequencies: {3: 30000, 4: 30000, 2: 30000, 1: 30000}
A few token frequencies: [('to', 106167), ('of', 71310), ('in', 64953), ('on', 47273), ('for', 36960)]
Label meanings: 1: World news, 2: Sports news, 3: Business news, 4: Sci/Tech news

Here's how we can get a sequence of tokens for a sentence with the cleaner, tokenizer, and vocabulary:

```

[54]: [vocab[token] for token in tokenizer(clean('Bangkok, or The Big Mango, is one_
↳of the great cities of Asia'))]

```

```

[54]: [3914, 96, 8, 291, 0, 11, 45, 7438, 1990, 3, 1057]

```

Let's make pipelines for processing a news story and a label:

```

[55]: text_pipeline = lambda x: [vocab[token] for token in tokenizer(clean(x))]
label_pipeline = lambda x: int(x) - 1

```

Here's how to create dataloaders for the training and test datasets:

```

[56]: import torch
from torch.utils.data import DataLoader
from torch.nn.utils.rnn import pad_sequence
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

def collate_batch(batch):
    label_list, text_list, length_list = [], [], []
    for (_label, _text) in batch:
        label_list.append(label_pipeline(_label))
        processed_text = torch.tensor(text_pipeline(_text), dtype=torch.int64)
        length_list.append(processed_text.shape[0])
        text_list.append(processed_text)
    label_list = torch.tensor(label_list, dtype=torch.int64)
    text_list = pad_sequence(text_list, padding_value=0)
    length_list = torch.tensor(length_list, dtype=torch.int64)
    return label_list.to(device), text_list.to(device), length_list.to(device)

train_iter = AG_NEWS(split='train')
train_dataset = list(train_iter)
test_iter = AG_NEWS(split='test')
test_dataset = list(test_iter)

```

```
train_dataloader = DataLoader(train_dataset, batch_size=8, shuffle=True,
    ↪collate_fn=collate_batch)
test_dataloader = DataLoader(test_dataset, batch_size=8, shuffle=False,
    ↪collate_fn=collate_batch)
```

Here's how to get a batch from one of these dataloaders. The first entry is a 1D tensor of labels for the batch (8 values between 0 and 3), then a 2D tensor representing the stories with dimension T x B (number of tokens x batch size).

```
[57]: batch = next(enumerate(train_dataloader))
print(batch)
```

```
(0, (tensor([3, 1, 0, 0, 0, 2, 3, 1], device='cuda:0'), tensor([[ 491,   136,
1145, 59356, 5454, 3052, 1388, 9830],
[ 3516,   895,   272,  1124,  1435,    2,  1352, 231994],
[ 4962,  2025,   424,  5895,  1637,   251,  3800,   242],
[ 50357,  2308,    16, 54052,  1083,  9010,   220,    6],
[  491,    23, 49994, 13942,   846,    6,   457,   22],
[  16,  5530,    19,    39,    72, 18237,    6,  2812],
[213605, 41516,    19, 12709,    72,    60, 96872,  9830],
[   65,  7562,   744, 39289,  1486,   177,    19, 140479],
[  201, 40739,   424,    25,    56,   17,    19, 227501],
[  7628,   540, 72416, 29607,  2217,  3052, 165581,   634],
[ 50357, 39511, 71440,   843,  1435,    84,   363,    4],
[ 29398, 11110,  1145,    72,  8047,  1913,   162,  1191],
[   129, 12040,   272,    72, 57174,    24,  18318,  9286],
[ 93032,  2744,    13,  1528,   532,   251,  1388,  2512],
[    8,   386,   42,    3,    6,  9010,   53,   54],
[ 1232,   38, 46270, 12709,   122,   222,    11, 86764],
[   77,    9,   306, 174853,    15,   81,  64921,  1106],
[   24,   11,   169, 25762,    56,    6,   317, 62271],
[   29, 3682,  1748,  1399, 12892, 18237, 11578,  1417],
[  841,  166,    23, 59356, 13762,    60,    77,  5311],
[179826,   11,  6132,    73, 10927,    4,  15087,  1309],
[   25,   539, 69627,  6815,  6858,   742,    4,    0],
[ 14337,    2,    35,  7726,  8903,  1229,   131,    0],
[ 46362,  207,  1015, 24279,    12,  2211,  7894,    0],
[244690,  540,  1304,   378,    34,   149,    3,    0],
[    0,   94, 32492, 192727,   840, 15387, 107799,    0],
[    0,   92,    3,   106,    2,  4192,  11001,    0],
[    0,  136,   23,    86,  4128,    78,    6,    0],
[    0,   895,  1779,  3912,    23, 20429,   67,    0],
[    0,  1020,    0,    13,  5288, 12909, 96871,    0],
[    0,  7984,    0,   213,   153,   751,   60,    0],
[    0,   866,    0, 135719, 88659, 128079,    4,    0],
[    0, 130247,    0,    26,    21,  3159, 27837,    0],
[    0, 41516,    0, 225089,   618,  2891,    0,    0],
[    0,  6464,    0,    4,    0,   996,    0,    0],
```

```

[ 0, 801, 0, 1555, 0, 227468, 0, 0],
[ 0, 11110, 0, 18, 0, 0, 0, 0],
[ 0, 18, 0, 253998, 0, 0, 0, 0],
[ 0, 41516, 0, 6391, 0, 0, 0, 0],
[ 0, 14744, 0, 152416, 0, 0, 0, 0],
[ 0, 609, 0, 1077, 0, 0, 0, 0],
[ 0, 0, 0, 10437, 0, 0, 0, 0],
[ 0, 0, 0, 843, 0, 0, 0, 0]],
device='cuda:0'), tensor([25, 41, 29, 43, 34, 36, 33, 21],
device='cuda:0'))

```

1.1 Question 1, 10 points

The vocabulary currently is too large for a simple one-hot embedding. Let's reduce the vocabulary size so that we can use one-hot. First, add a step that removes tokens from a list of "stop words" to the `text_pipeline` function. You probably want to remove punctuation (',', ',', '-', etc.) and articles ("a", "the").

Once you've removed stop words, modify the vocabulary to include only the most frequent 1000 tokens (including 0 for an unknown/infrequent word).

Write your revised code in the cell below and output the 999 top words with their frequencies:

```

[6]: # Place code for Question 1 here
def clean(line):
    line = line.replace('\\', ' ')
    #Add for question one
    line = re.sub('(\s+)(a|an|and|the)(\s+)', '\1\3', line)
    line = re.sub('[%s]' % re.escape(string.punctuation), '', line)
    return line

```

1.2 Question 2, 30 points

Next, let's build a simple RNN for classification of the AGNews dataset. Use a one-hot embedding of the vocabulary entries and the basic RNN from Lab 10. Use the lengths tensor (the third element in the batch returned by the dataloaders) to determine which output to apply the loss to.

Place your training code below, and plot the training and test accuracy as a function of epoch. Finally, output a confusion matrix for the test set.

Do not spend a lot of time on the training! A few minutes is enough. The point is to show that the model is learning, not to get the best possible performance.

```

[90]: # Place code for Question 2 here
import torch
import torch.nn as nn
import torch.optim as optim
num_class = len(set([label for (label, text) in train_iter]))

```

```

vocab_size = len(vocab)
emsize = 64
n_hidden = 128
n_tag = 4

class RNN(nn.Module):
    #Create a constructor
    def __init__(self, vocab_size, input_size, hidden_size, output_size):
        super(RNN, self).__init__()
        self.embedding = nn.EmbeddingBag(vocab_size, input_size, sparse=True)
        self.hidden_size = hidden_size
        self.rnn_cell = nn.RNN(input_size, hidden_size)
        self.h2o = nn.Linear(hidden_size, output_size)
        self.softmax = nn.LogSoftmax(dim = 1)

    #create a forward pass function
    def forward(self, input_, offsets_, hidden = None, batch_size = 1):
        embedded = self.embedding(input_, offsets_)
        out, hidden = self.rnn_cell(embedded, hidden)
        output = self.h2o(hidden.view(-1, self.hidden_size))
        output = self.softmax(output)
        return output, hidden

    def init_hidden(self, batch_size = 1):
        #function to init the hidden layers
        return torch.zeros(1, batch_size, self.hidden_size)

#train
net = RNN(vocab_size, emsize, n_hidden, num_class).to(device)
criterion = nn.NLLLoss()
opt = optim.SGD(net.parameters(), lr=0.01, momentum=0.9)

```

```

[91]: import time

def train(dataloader):
    net.train()
    total_acc, total_count = 0, 0
    log_interval = 500
    start_time = time.time()

    for idx, (label, text, offsets) in enumerate(dataloader):
        opt.zero_grad()
        predited_label = net(text, offsets)
        loss = criterion(predited_label, label)
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), 0.1)

```

```

    opt.step()
    total_acc += (predited_label.argmax(1) == label).sum().item()
    total_count += label.size(0)
    if idx % log_interval == 0 and idx > 0:
        elapsed = time.time() - start_time
        print('| epoch {:3d} | {:5d}/{:5d} batches | '
              '| accuracy {:.3f}|'.format(epoch, idx, len(dataloader),
                                          total_acc/total_count))

        total_acc, total_count = 0, 0
        start_time = time.time()

def evaluate(dataloader):
    net.eval()
    total_acc, total_count = 0, 0

    with torch.no_grad():
        for idx, (label, text, offsets) in enumerate(dataloader):
            predited_label = net(text, offsets)
            loss = criterion(predited_label, label)
            total_acc += (predited_label.argmax(1) == label).sum().item()
            total_count += label.size(0)
    return total_acc/total_count

```



```

[92]: EPOCHS = 10 # epoch
      LR = 5 # learning rate
      BATCH_SIZE = 64 # batch size for training

      for epoch in range(1, EPOCHS + 1):
          epoch_start_time = time.time()
          train(train_dataloader)
          accu_val = evaluate(valid_dataloader)
          if total_accu is not None and total_accu > accu_val:
              scheduler.step()
          else:
              total_accu = accu_val
          print('-' * 59)
          print('| end of epoch {:3d} | time: {:.5.2f}s | '
                '| valid accuracy {:.3f}|'.format(epoch,
                                                  time.time() - epoch_start_time,
                                                  accu_val))

          print('-' * 59)

```

```

ValueError                                Traceback (most recent call
↳ last)

<ipython-input-92-6cc7538ef1ad> in <module>
      5 for epoch in range(1, EPOCHS + 1):
      6     epoch_start_time = time.time()
----> 7     train(train_dataloader)
      8     accu_val = evaluate(valid_dataloader)
      9     if total_accu is not None and total_accu > accu_val:

<ipython-input-91-9686dc2f9267> in train(dataloader)
      9     for idx, (label, text, offsets) in enumerate(dataloader):
     10         opt.zero_grad()
---> 11         predited_label = net(text, offsets)
     12         loss = criterion(predited_label, label)
     13         loss.backward()

/usr/local/lib/python3.6/dist-packages/torch/nn/modules/module.py in
↳ _call_impl(self, *input, **kwargs)
     887         result = self._slow_forward(*input, **kwargs)
     888         else:
--> 889         result = self.forward(*input, **kwargs)
     890         for hook in itertools.chain(
     891             _global_forward_hooks.values(),

<ipython-input-90-0e6b9c63588a> in forward(self, input_, offsets_,
↳ hidden, batch_size)
     21     #create a forward pass function
     22     def forward(self, input_,offsets_,hidden = None, batch_size = 1):
---> 23         embedded = self.embedding(input_,offsets_)
     24         out, hidden = self.rnn_cell(embedded, hidden)
     25         output = self.h20(hidden.view(-1, self.hidden_size))

/usr/local/lib/python3.6/dist-packages/torch/nn/modules/module.py in
↳ _call_impl(self, *input, **kwargs)
     887         result = self._slow_forward(*input, **kwargs)
     888         else:
--> 889         result = self.forward(*input, **kwargs)
     890         for hook in itertools.chain(
     891             _global_forward_hooks.values(),

```

```

/usr/local/lib/python3.6/dist-packages/torch/nn/modules/sparse.py in
↳forward(self, input, offsets, per_sample_weights)
    340                 self.max_norm, self.norm_type,
    341                 self.scale_grad_by_freq, self.mode,
↳self.sparse,
    --> 342                 per_sample_weights, self.
↳include_last_offset)
    343
    344     def extra_repr(self) -> str:

/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py in
↳embedding_bag(input, weight, offsets, max_norm, norm_type, scale_grad_by_freq,
↳mode, sparse, per_sample_weights, include_last_offset)
    2045         ", as input is treated is a mini-batch of"
    2046         " fixed length sequences. However, found "
    -> 2047         "offsets of type {}".format(type_str)
    2048     )
    2049     offsets = torch.arange(0, input.numel(), input.size(1),
↳dtype=input.dtype, device=input.device)

```

```

ValueError: if input is 2D, then offsets has to be None, as input is
↳treated is a mini-batch of fixed length sequences. However, found offsets of
↳type <class 'torch.Tensor'>

```

1.3 Question 3, 10 points

Next, replace the SRNN from Question 2 with a single-layer LSTM. Give the same output (training and testing accuracy as a function of epoch, as well as confusion matrix for the test set). Comment on the differences you observe between the two models.

[8]: *# Place code for Question 3 here*

1.4 Question 4, 10 points

Explain how you could use the Transformer model to perform the same task you explored in Questions 2 and 3. How would attention be useful for this text classification task? Give a precise and detailed answer. Be sure to discuss what parts of the original Transformer you would use and what you would have to remove.

Answer: Transformer model can use a sequence of text as a input of model. This model can do as parallelization of sequence. For this task, there are many tag or many word in sentence. self-Attention machanism could be help for AgNews dataset which it try to looks at an input sequence and decides at each step which other parts of the sequence are important. To modify code for

self-Attention, we need to separate part into decoder and encoder, and add more linear layer as a V, Q, K (multihead attention) after that we use same softmax function for output.

1.5 Question 5, 10 points

In Lab 13, you implemented a DQN model for tic-tac-toe. Your method learned to play against a fairly dumb `expert_action` opponent, however. Also, DQN has proven to be less stable than other methods such as Double DQN, also discussed in Lab 13.

Explain below how you would apply double DQN and self-play to improve your tic-tac-toe agent. Provide pseudocode for the algorithm below.

Answer: Double DQN is part of RL model which contains 2 neural network models. First, model learns during the experience play as same as DQN. Second, copy last episode of the first model to compare for Q-value. If values of the second model are lower than the main model, we use the second model to attain Q-value. Sometime DQN overestimates the reward so double DQN decouples the actions selection from the action evaluation. To apply double DQN for tic-tac-toe, the input is the same as DQN but we need to create 2 neural networks. The first neural network will decide which one is the best next action and then the second network evaluates this action to know the Q-value.

```
[ ]: #Pseudo code
def select_greedy_actions(states: torch.Tensor, q_network: nn.Module) -> torch.
    ↪Tensor:

    _, actions = q_network(states).max(dim=1, keepdim=True)
    return actions

def evaluate_selected_actions(states: torch.Tensor,
                             actions: torch.Tensor,
                             rewards: torch.Tensor,
                             dones: torch.Tensor,
                             gamma: float,
                             q_network: nn.Module) -> torch.Tensor:
    next_q_values = q_network(states).gather(dim=1, index=actions)
    q_values = rewards + (gamma * next_q_values * (1 - dones))
    return q_values

def double_q_learning_update(states: torch.Tensor,
                             rewards: torch.Tensor,
                             dones: torch.Tensor,
                             gamma: float,
                             q_network_1: nn.Module,
                             q_network_2: nn.Module) -> torch.Tensor:

    actions = select_greedy_actions(states, q_network_1)
    q_values = evaluate_selected_actions(states, actions, rewards, dones, ↵
    ↪gamma, q_network_2)
    return q_values
```

2 Pseudo code

q_network1 is first neural network

q_network2 is secone neural network

Select_greedy_actions(state,q_network)

Selection action for q_network

evaluate_action(state,action,reward,done,q_network)

findding a next q value from q_network

evaluate q value

DDQN_update(state,reward,done,q_network1,q_network2)

selection from Select_greedy_actions by using state and q q_network1

evaluate action from evaluate_action which using q_network2 which return q value from secon

2.1 Question 6, 30 points

Based on your existing DQN implementation, implement the double DQN and self-play training method you just described. After some training (don't spend too much time on training -- again, we just want to see that the model can learn), show the result you playing a game against your learned agent.

```
[9]: # Code for training and playing goes here
```

```
[ ]:
```